

Stylized Facts about AI & Human Capabilities

Tom Cunningham & Ali Merali | 2026-05-17

We discuss stylized facts about the capabilities of recent AI models and how they compare to human capabilities. We particularly draw attention to a common latent factor of capabilities among AI models, and the fact that this latent factor is different from the distribution of human capabilities. We discuss various candidates for the “missing capability” which would unlock a greater economic impact.

Table of contents

1	Overview	3
1.1	We wish to characterize AI capabilities	3
1.2	There is no standard measure of machine intelligence	5
1.3	There is no standard taxonomy of AI capabilities.	6
1.4	There does appear to be a single latent component of intelligence across LLMs	12
1.5	Latent LLM intelligence appears to be distinct from latent human intelligence.	14
1.6	There is <i>some</i> bottleneck on LLM abilities, but disagreement about where	16
1.7	There are many candidates for the bottleneck . .	17
2	LLM Performance on Tasks	24
2.1	The time between benchmark introduction and saturation has been decreasing	24
2.2	LLM performance is highly correlated across benchmarks	25
2.3	The best predictor of benchmark performance is training compute.	27
2.4	The effective scale of frontier models has been growing at around 15X per year.	29

2.5	Benchmark performance significantly improves with test-time scaling on some tasks.	29
2.6	Benchmark tasks that are harder for humans are typically harder for LLMs.	30
2.7	LLM performance degrades significantly with context length.	30
2.8	LLM performance on benchmarks is highly variable within the same task.	31
2.9	LLM Performance on related sub-tasks is highly variable.	31
2.10	LLM performance is sensitive to prompt wording.	31
2.11	Benchmark tasks that take humans longer are typically harder for LLMs.	32
2.12	LLMs perform particularly well on math and programming competition problems.	33
2.13	LLMs show some ability to autonomously complete software engineering tasks.	33
2.14	LLMs are showing ability to improve on algorithmic tasks.	33
2.15	LLMs are able to autonomously complete real-world tasks with significant economic value. . . .	33
2.16	There are few benchmarks in which humans significantly outperform computers.	34
2.17	LLMs show limited ability to perform actions in the real world.	34
2.18	LLMs show limited ability to play single-player and multiplayer games.	35
2.19	LLMs hallucinate with obscure or long-form answers.	35
2.20	LLMs struggle to give calibrated answers.	36
3	LLM Augmentation Studies	36
3.1	The productivity effects of LLM augmentation varies widely	37
3.2	There is some evidence of larger effects on lower-productivity workers	39
3.3	There are few benchmarks of augmentation	39
3.4	Human + AI teams do not reliably outperform the better of human or AI performance individually	39
3.5	Humans often misperceive the abilities of LLMs.	40

4	Implications for LLM Economic Impact	41
4.1	The price of inference for a fixed level of capability has been falling at 10X/year or faster. . . .	41
4.2	The cost of inference for frontier capabilities has remained roughly stable.	42
4.3	Inference costs of serving a model are now roughly proportional to the initial training costs.	43
4.4	LLMs are almost always cheaper than humans .	43
4.5	The economic impact of LLMs has been modest so far	44
4.6	We may not be able to direct technical change .	45
4.7	A univariate model of machine and human intelligence will be a poor fit	46
4.8	Wages are unlikely to be pinned down by the price of compute	47
5	Appendix / Offcuts	48
	References	48

Thanks to many for their suggestions, especially those from the Economic Research team at OpenAI, and Boaz Barak. The document includes significant suggestions from OpenAI staff¹ (Pamela Mishkin, Zoe Hitzig, Gawesha Weeratunga, Ronnie Chatterji), and external academics: Catherine Tucker, David Deming, Erik Brynjolfsson, Anton Korinek, Pascual Restrepo.

1 Overview

1.1 We wish to characterize AI capabilities

Our goal is to give a crisp characterization of AI capabilities.

We wish to answer questions like these:

1. What types of tasks can AI models do / cannot do?
2. How do abilities differ among different AI models?
3. How do AI abilities differ from human abilities?
4. What are reasonable expectations about future AI abilities?

A running theme will be that the characterization of AI capabilities is *difficult*.

There is a long history of struggle in making accurate

generalizations about the capabilities of artificial intelligence. There is little consensus on how to describe the abilities of contemporary models.

In this note we list a series of well-documented facts about AI capabilities, and make some more opinionated generalizations. The target audience is economists who are interested in modeling the effects of AI on various situations.

It is useful to think of a matrix of agents and tasks.

Suppose we can observe whether an agent can perform a task, for a long list of agents and tasks. The agent could be a human, a computer, or a hybrid. We can represent each observation in a matrix with agents and columns and tasks as rows:

	Alice	Bob	Charlie	AI in 2020	AI in 2024	AI in 2028
2 Basic Algebra Questions						?
Trivial Pursuit Questions						?
Pass LSAT						?
Complete Super Mario Bros						?
Prove Reimann hypothesis						?

We wish to find a low-dimensional factorization of this matrix.

In practice the matrix is enormous. We can think of the capability problem as finding a low-dimensional representation of this matrix, i.e. some small set of general-purpose capabilities that are highly predictive of whether an agent can do a task.¹

A single-dimensional factorization would mean each agent would have a latent ability, and each task would have a latent difficulty. However we will see that

¹We can think of the reduction either as Principal Components Analysis (we take weighted averages of all rows), or Principal Feature Analysis (we take a subset of rows that are highly predictive).

a single-factor is often a poor fit, i.e. the relatively difficulty of tasks varies substantially across agents.

1.2 There is no standard measure of machine intelligence

Many proposed measures of AGI have already been satisfied.

Scott Alexander [says](#) “*The history of AI is people saying ‘We’ll believe AI is Actually Intelligent when it does X!’ – and then, after AI does X, not believing it’s Actually Intelligent.*”

Some examples:

Measure	Year passed
Chess – “ <i>Chess is generally considered to require ‘thinking’ for skilful play.</i> ” — Shannon, 1950.	1997
Winograd Schemas - “ <i>when people pass the test, we would say they were thinking.</i> ”	2019
MMMU – “ <i>Benchmark for Expert AGI.</i> ”	2025
ARC-AGI – “ <i>the only formal benchmark of AGI progress.</i> ”	2024
AGIEval - “ <i>general abilities ... in tasks pertinent to human cognition and problem-solving.</i> ”	2025

There are some proposed definitions but none have become canonical:

- Legg and Hutter (2007b) proposed “intelligence measures an agent’s ability to achieve goals in a wide range of environments.”²

²Additionally, Legg and Hutter (2007a) surveyed 70 different definitions of intelligence as of 2007.

- Chollet (2019) proposed intelligence be defined as “skill-acquisition efficiency”, and proposed a formal benchmark (ARC-AGI).
- Morris et al. (2023) proposed five “levels of AGI”, based on the percentile human performance across all tasks. E.g. they classified ChatGPT as being able to do a *“wide range of non-physical tasks, including metacognitive tasks like learning new skills equal to or somewhat better than an unskilled human.”*
- In 2024 OpenAI internally developed a five-level scale of AGI, but it has not been publicly released.³
- Hendrycks et al. (2025) proposed an AGI definition based on 10 different sub-dimensions.

1.3 There is no standard taxonomy of AI capabilities.

There are hundreds of different AI benchmarks of capability.

These benchmarks are often informally grouped into types which test different types of ability, e.g. “knowledge”, “reasoning”, “programming”, “agentic ability”, however there is no generally-accepted taxonomy.

³<https://www.bloomberg.com/news/articles/2024-07-11/openai-sets-levels-to-track-progress-toward-superintelligent-ai>




Figure 8: Ability profiles of the 15 LLMs. An ability of l means that there is 50% probability of the model to succeed on questions at demand level l (that leads to some abilities being beyond 5). In contrast to radial plots usually shown for LLMs in the literature [7, 51] the values shown here are actual abilities on a ratio scale $(0, \infty)$ and the values (in expectation) are more robust to changes in the difficulty distribution of the benchmarks used. In Figure 10, we additionally show the ability profiles for the 15 LLMs using the ADeLe-Light battery v.1.0, which yields nearly the same profiles. In Appendix 9.2, we show clear scaling curves of model abilities as a function of #parameters.

Figure 1: 18-dimensional taxonomy from Lexin Zhou et al. (2025)

Table 2: Model performance on Bloom’s Taxonomy’s Cognitive Dimension and Knowledge Types. Model scores marked with * are based only on our own evaluation on BBH and AGIEval.




Figure 2: Performance of models mapped to Bloom’s Taxonomy. Model scores marked with * are based only on BBH and AGIEval. Unmarked scores are based on all benchmarks.

Figure 2: 10-dimensional taxonomy from Huber and Niklaus (2025)




Figure 1: The capabilities of GPT-4 and GPT-5. Here GPT-5 answers questions in ‘Auto’ mode.

Figure 3: 10-dimensional taxonomy from Hendrycks et al. (2025)




Figure 4: From Ilić and Gignac (2024), showing their preferred 2-factor model




Figure 4: Projection of new models

Figure 5: From Maimon et al. (2025)

Some taxonomies of AI capabilities have been proposed but none have been widely adopted.

Burnell et al. (2023) look at variation across 29 LLMs and 27 benchmarks (all prior to GPT-4). The first 3 factors explain 33%, 31%, and 17% of performance, and they interpret those factors as “reasoning”, “comprehension” and “core language modeling.” They say “capabilities are not monolithic. Instead, they are better explained by three well-delineated factors that represent reasoning, comprehension and core language modeling.”

Ilić and Gignac (2024) tested 591 LLMs on 12 different evaluations, as of March 2024. In their preferred specification the first factor had a loading of 0.81, implying an R^2 of 0.66. They also found that the “number of LLM parameters correlated positively with this general ability.”

Polo et al. (2025) propose a decomposition with three latent factors: reasoning, knowledge, and instruction-following. However they do not seem to report overall variance explained by each of the factors.

Maimon et al. (2025) analyze 60 LLMs on 44 tasks, and propose eight latent factors.

Lexin Zhou et al. (2025) propose an index of 18 different human-interpretable capabilities (ADeLe). They use an LLM to label 16,000 “task instances” from 20 benchmarks, using rubrics to define each capability. This allows them to predict LLM performance on a task just based on its content, without any prior LLM successes. The scores on each dimension are 1-5, they note that they scores approximately correspond to frequency of achievement among humans, where score L will be achieved by $1/10^{L-1}$ people.⁴ response.”)

OECD (2025) organize capabilities across nine dimensions. The criteria are mostly qualitative, and their scoring is done by a panel of experts with reference to various benchmark scores. OECD’s dimension: Language; Social interaction; Problem solving; Creativity; Metacognition

⁴“an item is at level l if l is the highest number such that, in at least 95% of samples of $n = 10^l$ individuals, there is at least one correct.”

and critical thinking; Knowledge, learning and memory; Vision; Manipulation; Robotic intelligence.

Huber and Niklaus (2025) organize benchmark tasks according to “Bloom’s Taxonomy”, a hierarchical framework designed to classify educational learning objectives for humans, with 6 levels: Remember, Understand, Apply, Analyze, Evaluate, Create. They say there are a lack of benchmarks to measure the two highest levels (“evaluate” and “create”).

Hendrycks et al. (2025) define capabilities along 10 dimensions, based on the Cattell-Horn-Carroll (CHC) theory of cognitive abilities. They give testable evals for only some of these capacities.

Many taxonomies are based on components of between-agent variation.

Many taxonomies are based on latent factors in variation in performance in a range of tasks.

Some taxonomies are based on variation among human abilities (e.g. Hendrycks et al. (2025), Huber and Niklaus (2025)). Some taxonomies are based on variation across LLM abilities (Burnell et al. (2023), Ilić and Gignac (2024), Polo et al. (2025), Maimon et al. (2025), Lexin Zhou et al. (2025)).

A difficulty with both of these approaches is that they are designed to measure variation *within* one type of agent (LLM or human), not *between* humans and LLMs. There is reason to expect that there exist important human-LLM differences that are orthogonal to within-human or within-LLM variation, e.g. types of tasks where humans show little variation, and LLMs show little variation, but there are large human-LLM differences in ability.⁵


1.4 There does appears to be a single latent component of intelligence across LLMs

LLM abilities are highly correlated across models.

LLM performance is very highly correlated across benchmarks: if model A beats model B on benchmark X, then it very likely also beats model B on benchmark Y.

We discuss some literature below which tries to quantify the strength of this correlation.

⁵A similar argument is made by Zachary Tidler et al. [here](#).



(a) PCA explained variance

Figure 6: From Ruan, Maddison, and Hashimoto (2024)

AI researchers often treat intelligence as a scalar.

A common sentiment among AI researchers is that it is best to spend time on making a model generally better, rather than fine-tuning for a specific case.

This sentiment famously named the “bitter lesson” by Sutton (2019):

“general methods that leverage computation are ultimately the most effective, and by a large margin.”

Much work on “scaling laws” treats model intelligence as a scalar.

A large part of academic and industry work on AI is spent

estimating “scaling laws”, i.e. predicting model performance based on various inputs (e.g. training compute, training data size, model scale, and test-time compute). Most work on scaling laws uses a scalar outcome to represent model intelligence, e.g. prediction error on a corpus, or performance on a difficult benchmark. In principle scaling laws could estimate a *vector* of intelligence outputs however this is relatively rare, researchers tend to treat a single scalar as a sufficient proxy for model intelligence.⁶

Some theories of the latent component.

- The Bitter Lesson.
 - The “cognitive core”: (Sean Goedeck)
 - The Manifold Hypothesis.
 - The Platonic Representation hypothesis.

1.5 Latent LLM intelligence appears to be distinct from latent human intelligence.

Humans have a fairly positive correlation among abilities.

There is a reasonable correlation among different measures of general ability (IQ, g-factor).

⁶polo2025sloth propose a decomposition with three latent factors: reasoning, knowledge, and instruction-following.





Figure 7: Illustration of LLM vs humans with distribution of ability.



There is a reasonable correlation between years of education, IQ, standardized tests score, and wage.

Ideally we would compare human and LLM performance on the same tasks.

If we had a full matrix of human and LLM performance we would be able to compare latent factors of each. Unfortunately there are few datasets which collect human performance across a range of distinct benchmarks.

We can compare the average human with the average LLM.

Many benchmarks do report performance for the average or expert human, so we can make those comparisons. Additionally we can run LLMs on many human standardized tests and exams.

Latent factors for human and LLM ability appear to be significantly different.

Aptitude tests that are predictive for humans are not predictive for LLMs. GPT-4 was able to get human-level scores at most aptitude tests: SAT, LSAT, civil service. But GPT-4 did not have the aptitudes passing such tests implies for humans.

LLMs are able to do tasks no human can: LLMs can translate between pairs of languages which no human knows (Aymara to Kazakh), can combine knowledge from two very different domains.

1.6 There is *some* bottleneck on LLM abilities, but disagreement about where

In 2023 many people thought GPT-4 could do a significant fraction of work:

- Eloundou et al. (2023) estimated that GPT-4 would halve the time required to perform a task for between 14% and 56% of all tasks in the economy.
- Chui et al. (2023) (McKinsey) claimed “*the total percentage of hours that could theoretically be automated by integrating technologies that exist today ... [is] 60–70 percent.*”

In retrospect these estimates seem to have been overly optimistic. There is some missing capability in GPT-4 level ability.

AI researchers are not sure [Andrej Karpathy](#), Nov 2024: what is the bottleneck.

“Even though by many accounts (evals), LLMs are inching well into top expert territory (e.g. in math and coding etc.), you wouldn’t hire them over a person for the most menial jobs. It’s an interesting challenge how we can create evals for all the “easy” stuff that is secretly hard. Very

long context windows,
coherence, autonomy,
common sense, multi-
modal I/O that works,
... How do we build
good “menial job”
evals?”*

Shunyu Yao, April 2025:

“AI has beat world
champions at chess
and Go, surpassed
most humans on SAT
and bar exams, and
reached gold medal
level on IOI and IMO.
But the world hasn’t
changed much, at least
judged by economics
and GDP. I call this
the utility problem,
and deem it the most
important problem for
AI.”*

1.7 There are many candidates for the bottleneck

Here we list a variety of abilities that LLMs lack, which might explain the bottleneck to adoption. To make the bottleneck question concrete: LLMs can pass many tests used to test for human aptitudes (SAT, LSAT, civil service exams), and many interview questions, yet LLMs clearly lack the aptitude to perform the associated jobs. Thus there must be some other missing skill.

Capability	Benchmark	Apr 2023 (GPT-4)	Dec 2025	Human
Deductive reasoning	MATH ⁷	23%	98%	
	GPQA ⁸	36%	93%	
Information-gathering	GAIA[[^] GAIA]	15%	88%	
Accuracy / Hallucination	SimpleQA	~30%	72%	
Long context	Verified ⁹			
	Needle ¹⁰	50% @ 128K	99.7% (1M?)	
Multimodal inputs	MMMU-Pro ¹¹	~35%	81%	85%
In-context learning	ARC-AGI ¹²	7%	88%	77-98%
	ARC-AGI 2 ¹³	~0%	45%	98%

Bottleneck	Benchmarks	Benchmarks
Continual learning.	Amodei, Karpathy, Hassabis.	
Novel situations	RL generalization, world models.	
Interactive decision-making.		BALROG, 44%

⁷

⁸

⁹<https://arxiv.org/pdf/2509.07968v1> and <https://blog.google/products/gemini/gemini-3/#gemini-3>

¹⁰<https://cloud.google.com/blog/products/ai-machine-learning/the-needle-in-the-haystack-test-and-how-gemini-pro-solves-it>

¹¹<https://mmmu-benchmark.github.io/#leaderboard> and <https://blog.google/products/gemini/gemini-3/#gemini-3>

¹²<https://arcprize.org/leaderboard>

¹³<https://arcprize.org/leaderboard>

Data density.

Werra (2025) argues that jaggedness is mainly due to availability of data.

Accuracy.

LLMs have often made confident false claims (“hallucinations”), and this has clearly been a barrier to their adoption. This behavior is not penalized in classic benchmarks which score only right and wrong answers, and so the LLM has an incentive to guess. Similarly user pairwise preference does not penalize hallucinations when the user is not immediately aware of the accuracy of claims.

Benchmarks can test for accuracy in a variety of ways: (1) score based on confidence; (2) elicit open-ended responses, with a penalty based on claims that are false or misleading.

Deductive reasoning.

Many questions involving reasoning are drawn from a fairly stylized pool, and LLMs historically performed poorly on questions that are novel, or have misleading associations (“Alice in Wonderland”). Since the launch of reasoning models in late 2024 LLMs have become far better at deductive problems.

Novel situations.

LLMs sometimes fail on problems that are very dissimilar to

their training data (“out of distribution”). Because LLMs are trained on enormous amounts of data their skills could reflect memorization or interpolation rather than general intelligence.

Vafa et al. (2024) argue that Transformer-based LLMs can generalize well along some axes but badly along others. They say “such incoherence creates fragility: using a generative model to solve related but subtly different tasks can lead to failures.” They propose tests of world model coherence.

High context.

Questions on aptitude tests have short prompts, while real-world problems involve rich context.

Continual learning.

Typically LLMs have been used in a *stateless* manner, where the model weights are fixed and only the prompt changes. This means that any learning about the environment must happen with a separate system, or by through recursive updates to the prompt. This has been cited by many people as a critical bottleneck: Karpathy, Hassabis, Altman, Amodei.

[Hassabis, August 2025](#)

[Dwarkesh Patel, July 2025:](#)

“the fundamental
problem is that LLMs

don't get better over time the way a human would. The lack of continual learning is a huge huge problem."

Dario Amodei, August 2025:

"Models learn within the context ... Maybe we'll train the model in such a way that it is specialized for learning over the context. You could, even during the context, update the model's weights ... So, there are lots of ideas that are very close to the ideas we have now that could perhaps do [continual learning]."

Noam Brown, Nov 2025:

"More research breakthroughs are probably needed to achieve AGI/ASI. (Continual learning and sample efficiency are two examples that researchers commonly point to.)"

Shane Legg:

"So there are all sorts of things like this [continual learning and visual reasoning], that we currently see. I don't think there are fundamental blockers on any of these things. ... So my expectation is over a number of years these things will all get addressed. But they're not there yet." "

Planning / long-horizon decision-making.

Dorn et al. (2025) argue that current LLMs are poor at “sequential coherence”, they estimate the time-length of ONET tasks as a proxy for sequential coherence required.

In-context learning.

Chollet (2019) argues that most benchmarks test for *existing* skills instead of requiring the agent to learn *new* skills, which is distinctive for human intelligence. He introduced the Abstraction and Reasoning Corpus (ARC) task to test for the ability to do skill acquisition (also sometimes called in-context learning).

LLMs made significant progress on the ARC benchmark, approaching human levels in 2024. However a new version was introduced (ARC-AGI-2), which is again relatively straight-forward for a human but difficult for an LLM.

Multimodal inputs and outputs.

LLMs primarily operate on text inputs and outputs.

Interactive decision-making.

Questions on aptitude tests do not require a series of decisions, navigating an uncertain situation.

Predictability.

Some people have claimed a

bottleneck is the *unpredictability* of LLM abilities. This can be interpreted in two ways: (1) humans have systematically biased beliefs about LLM abilities; (2) i.e. even if LLM accuracy might be high over a broad class of questions, the accuracy varies across subclasses. Vafa, Rambachan, and Mullainathan (2024) say “when the cost of mistakes is high, more capable models can perform worse on the instances people choose to use them for because they are not aligned with the human generalization function.”

A related point is that LLMs can be inconsistent, giving qualitatively different answers to a question when the wording is slightly different.

Cost. (…)

I have not added separate category for “reliability,” meaning lack of randomness, because it is relatively to run an LLM in a mostly deterministic fashion, and so eliminate concerns about randomness.

Other discussions of bottlenecks:

- Toner (2025) discusses reasons for jaggedness: (1) easier vs harder to verify tasks; (2) can you put the context in the context window; (3) level of adversarialness; (4) cognitive vs physical.
- Heitmann et al. (2025)
 - Performance on tasks that are hard to verify
 - Performance on long tasks (that take human workers a long time)

- Performance in complex environments
- Sufficiently low error rate
- Meta-awareness and calibration
- Adaptability to the deployment environment
- Continual learning post-deployment
- Original insights


2 LLM Performance on Tasks

In this section we discuss performance on tasks that have a measurable output.

2.1 The time between benchmark introduction and saturation has been decreasing

If we had a cardinal scale of computer intelligence we could talk about the rate of improvement in a precise way. The best proxy for velocity is the history of AI benchmarks: new benchmarks are introduced when computers achieve human-level performance on the old benchmarks (“saturation”).

Since 2012 the rate of benchmark saturation seems to be increasing:



2.2 LLM performance is highly correlated across benchmarks





Fig. 3. Open LLM Leaderboard correlation matrix

Figure 6: Pairwise Correlations between benchmarks listed in LLMsSys.

Figure 9: From Spangher et al. (2024)



(a) PCA explained variance

Figure 10: From Ruan, Maddison, and Hashimoto (2024)

Frontier models that surpass previous models on one benchmark usually outperform them on most other benchmarks.

Spangher et al. (2024) construct an aggregation across benchmark scores (MPG = Model Performance and Goodness), which they find is highly correlated with Elo rating on LMArena. The strongest model in their dataset is **o1-preview**.

Ruan, Maddison, and Hashimoto (2024) create “observational scaling laws” from the performance of 77 LLMs (as of May 2024) on 8 benchmarks. They find the first component explains 80% of the variation, also that training compute is highly predictive of this factor (for open models,

which have known scale). They also fit scaling laws for the effect of inference-time techniques such as Chain-of-Thought and Self-Consistency.

Kipnis et al. (2025) find a small set of questions which are highly predictive of overall model scores, and that “a single underlying common factor whose Spearman correlation with the total score is $r = 0.93$.”

2.3 The best predictor of benchmark performance is training compute.

There is a well-known literature on “scaling laws” finding that pretraining loss (i.e. error in predicting the next token) is a predictable function of the training parameters.

- Kaplan et al. (2020) predict LLM performance in next-token prediction error, as a function of model size, dataset size, and training compute. They find scaling laws that are reliable across 7 orders of magnitude. Their equations implied that model size was the most important parameter (not dataset size).¹⁴
- J. Hoffmann et al. (2022) revisit scaling laws with more data and a different functional form:

$$\hat{\underset{\text{loss}}{L}}(\underset{\text{parameters}}{N}, \underset{\text{data}}{D}) \simeq \underset{\text{irreducible loss}}{E} + aN^{-.3} + bD^{-.3}$$

Their fit finds that training data and model size should optimally scale at the same rate. Their model predicted that it would be optimal to train a model with much

¹⁴“Larger models are significantly more sample-efficient, such that optimally compute-efficient training involves training very large models on a relatively modest amount of data and stopping significantly before convergence.”

smaller size and larger amount of data, compared to existing models. They trained such a model (“chinchilla”) and indeed found that it significantly outperformed other models with the same compute inputs.

- Choshen, Zhang, and Andreas (2024) surveys scaling-law best-practices in 2025, and largely confirms the functional-form and parameters proposed in J. Hoffmann et al. (2022).

The classical “scaling laws” literature predicts the fit on training loss, meaning next-token prediction error. There is a smaller literature which predicts performance on task benchmarks (note that these are usually from a model with post-training):




Figure 1: Aggregate benchmark performance is decently predictable from scaling-estimated loss, i.e., scaling of model size and dataset size. Loss, on the x-axis, is expressed as scaled compute equivalent per Section 2.1. Sigmoid fits are shown in black for both plots, with solid lines showing data included in fitting and dashed lines showing held-out extrapolation. 90% confidence intervals for fits are shown.

Figure 11: From Owen (2024)

- Owen (2024) find that model scale is highly predictive of benchmark performance, and the relationship remains relatively stable across five orders of magnitude.¹⁵ The strongest model in their dataset is GPT-4. This paper also discusses a number of previous contributions on benchmark scaling.
- Ruan, Maddison, and Hashimoto (2024) find tight relationships between training compute, simple capability

¹⁵“when extrapolating BIG-Bench Hard performance across one order of magnitude in compute, we observe average absolute errors of 6 percentage points.”

measures, and complex downstream metrics.

2.4 The effective scale of frontier models has been growing at around 15X per year.

The effective scale of frontier models takes into account both the increase in training compute used as well as the algorithmic efficiency in utilizing each unit of compute.

- [Epoch AI](#) (Jan 2025) estimate that over 2015-2025:
 1. Training compute on frontier LLMs has increased around 5X per year;
 2. Algorithmic efficiency increased around 3X per year.
 - Combined, these imply that the “effective” compute used on frontier models is growing at around 15X per year.

2.5 Benchmark performance significantly improves with test-time scaling on some tasks.

Test-time scaling (AKA inference-time compute) can be achieved in two ways:

1. Aggregating multiple models calls (self-consistency AKA majority vote Wang et al. (2023), tree-of-thought, etc.)
2. Training a model with a variable budget of inference-time tokens (“reasoning” models).

The returns to inference-time compute seem to vary by the nature of the problem.

Test-time scaling is not a perfect substitute for model scale: Wu et al. (2025) shows that small models with inference-time compute have performance that asymptotes below those of large models.

- Beeching, Tunstall, and Rush (n.d.) find that traditional test-time scaling (majority vote, best-of-N, beam search) shows sub-log-linear returns on a math problems benchmark.

- You (2025) (Epoch) say that existing evidence from OpenAI and DeepSeek shows log-linear relationships between test-time compute and performance on benchmarks, and say “this suggests that performance scales with reasoning training in a way similar to pre-training, at least for math and coding tasks.”
- Balachandran et al. (2025) compare inference-time compute across a variety of tasks and say “*Although inference-time scaling improves performance, its effectiveness varies between domains and tasks, with diminishing returns as task complexity increases.*”

2.6 Benchmark tasks that are harder for humans are typically harder for LLMs.

There is typically a positive correlation between the tasks that humans find more difficult and ones LLMs struggle with the most.

- Lucy Zhou et al. (2024) find a positive correlation between human difficulty and AI difficulty.
- Dreyfuss and Raux (2025) find no correlation between human and AI success rate across a set of math questions, however note that the results are using GPT-3.5.

2.7 LLM performance degrades significantly with context length.

As of July 2025, frontier models have an effective context length (where performance is >85% of a short-context task) of around ~16,000 words in the [NoLiMa Benchmark](#).

- LongBench (Bai et al. (2024)) is a benchmark for long-context understanding, found that the best models as of Sep 2024 (o1-preview) outperformed human baseliners given 15 minutes by 15% on short tasks (<32k context), this dropped to only a 5% gain for long-context tasks (>128k).

2.8 LLM performance on benchmarks is highly variable within the same task.

LLMs have stochastic outputs and they often have varying success rates when the same task is repeated. As a consequence an LLM’s average score on a benchmark can be very different from their best-of-n and worst-of-n scores (note that if output was deterministic then all three scores would be identical).

- Renze and Guven (2024) find “changes in temperature from 0.0 to 1.0 do not have a statistically significant impact on LLM performance for problem-solving tasks”
- Song et al. (2024) say “we observe that greedy decoding [temperature=0] generally outperforms sampling methods [temperature>0] for most evaluated tasks.”

2.9 LLM Performance on related sub-tasks is highly variable.

LLMs generally have much smaller correlations in their ability to successfully complete similar tasks than humans.

- Dreyfuss and Raux (2025) find that, for a given aggregate SAT score, humans have much larger variations across subjects than LLMs.

2.10 LLM performance is sensitive to prompt wording.

The performance of LLMs on certain benchmarks can vary significantly depending on the prompting given to the model.



- Mirzadeh et al. (2024) find that adding irrelevant instructions to a prompt significantly lowers performance on grade-school level questions. However [Andrew Mayne \(2024\)](#) shows that this performance penalty disappears when the prompt included a warning (“*this might be a trick question designed to confuse LLMs with additional information.*”).

- Valmeekam, Stechly, and Kambhampati (2024) finds that translating an ordinary problem into an artificial language significantly lowers performance on planning problems (e.g. “mystery blocksworld” in PlanBench). However o1-preview was still able to perform at 50%.
- Li et al. (2025) find that performance on math problems degrades somewhat when reworded. The declines are particularly severe for small open-source models.

2.11 Benchmark tasks that take humans longer are typically harder for LLMs.

There is a very strong negative correlation between how long it takes for a human to complete a task (in logarithmic terms) and the probability that an LLM will be able to complete it successfully.

- Kwa et al. (2025) find that the longest length of task that a frontier LLM has been able to complete with greater than 50% probability has been doubling every seven months over the last six years across a range of software engineering, cybersecurity, and general reasoning tasks. As of June 2025 it stands at just under two hours. A July 2025 update from METR shows that the time horizon is similar across other benchmarks testing quantitative and analytical skills and whilst there are larger differences in other domains (eg. web-browsing tasks) the rate of change over time is similar. The relationship between model success probability and the length of time it takes a human to complete a task is well fit with a logistic curve. The correlation in the errors of different models is high (79% between claude 3.7 and o1) even after accounting for task length and average model performance suggesting there are features of tasks beyond task length with predictive power for model performance.
- Ord (2025) argues that LLMs performance degradation over time is roughly consistent with a hypothesis that long tasks are random combinations of subtasks, each with a constant half-life of success.



2.12 LLMs perform particularly well on math and programming competition problems.

- Frontier LLMs can outperform the vast majority of competitors in math competitions (eg. [AIME](#), [MATH](#)). In a prominent programming competition, [Codeforces](#), o3's performance would place it as the 175th ranked competitive programmer in the world.

2.13 LLMs show some ability to autonomously complete software engineering tasks.

- Frontier models can autonomously resolve 75% of a set of curated issues in software systems on github ([SWE-bench-verified](#)).
- Frontier models can autonomously produce predictive models on 17% of Kaggle datasets with the same error as a highly skilled data scientist (bronze on [MLE-bench](#)).

2.14 LLMs are showing ability to improve on algorithmic tasks.

- Google (2025) reports a 0.7% saving in Google's world-wide compute resources) and aid in the design of new TPU chips.

2.15 LLMs are able to autonomously complete real-world tasks with significant economic value.

Frontier models are able to autonomously complete tasks with high market value. For instance, an open-source benchmark published by OpenAI finds:

- Miserendino et al. (2025) (SWE-lancer) measures LLM's abilities to autonomously complete software engineering and management tasks posted on UpWork priced from

\$50-\$32,000. As of May 2025, frontier LLMs were able to complete 40% of tasks (i.e. earn \$400,000 out of \$1M).

2.16 There are few benchmarks in which humans significantly outperform computers.


It’s challenging to find benchmarks in which humans can still outperform computers. A recent list of “[useful benchmarks that have human scores beyond AI SOTA](#)” finds benchmarks that require UX usage have so far (Nov 2024) resisted saturation.

Even on UX usage benchmarks, the difficulty for LLMs appear to be more about tooling (which may be more easily rectifiable) than planning capabilities (the OSWorld paper notes that 75% of failed attempts are due to inaccurate mouse clicks rather than planning mistakes).

2.17 LLMs show limited ability to perform actions in the real world.

LLMs are able to perform some simple discrete multi-turn tasks in simulated computer environments but have yet to saturate benchmarks of simple web-browsing tasks such as [Tau-Bench](#) and [OS-world](#).

- Anthropic reported on a [2025 project](#) for Claude to run as a “shopkeeper” for a vending machine inside the Anthropic offices. They found that the agent hallucinated payment accounts, sold items at a loss, and got talked into severe discounts.
- Vending-Bench (Backlund and Petersson (2025)) is a fully simulated vending machine setup: frontier agents routinely drive the business to bankruptcy within a few in-game weeks, suggesting a lack of long-term navigation abilities.




2.18 LLMs show limited ability to play single-player and multiplayer games.

Frontier LLMs struggle playing games that are simple for children to solve.

- Paul Calcraft [discusses](#) the state of LLM-gameplaying in early 2025: their performance in tic-tac-toe, connect 4, codenames, and chess, is still well below that of human experts.
- Payne and Alloui-Cros (2025) find that LLMs are “highly competitive” with the most successful strategies in playing iterated prisoners dilemma games.
- LMGame-Bench (Hu et al. (2025)) finds that as of early 2025, LLMs are still below human performance across 6 computergames: Sokoban, Super Mario Bros, Tetris, 2048, Candy Crush, and Ace Attorney.
- TextQuests (Phan et al. (2025)) find as of August 2025 LLMs are unable to solve text-based adventure games that would take human players around 30 hours and hundreds of actions to solve without assistance. With clues granted recent models are able to make greater progress-whilst no model prior to 2025 is able to solve any adventure completely, GPT-5 is able to solve 5/25.

9



2.19 LLMs hallucinate with obscure or long-form answers.

LLMs struggle with reliability and often make up plausible-sounding but factually incorrect hallucinations.

- GPT-4o and o1-preview hallucinate 30-60 % of the time when asked obscure factual questions, e.g. “Which Dutch player scored an open-play goal in the 2022 Netherlands vs Argentina game in the men’s FIFA World Cup?” ¹⁶.

¹⁶https://assets.ctfassets.net/kftzwdyauwt9/67qJD51Aur3eIc96iOfeOP/71551c3d223cd97e591aa89567306912/o1_system_card.pdf

- When evaluated on a proprietary benchmark of open-ended questions they included incorrect facts in their answers around 80 % of the time ¹⁷.

However, there has been recent progress made in reducing hallucination rates. GPT-5, released in August 2025, had six times fewer hallucinations than o3 as measured by LongFact -Wei et al. (2024) and FActScore -Min et al. (2023).

2.20 LLMs struggle to give calibrated answers.


LLMs usually are not aware of what they know or do not know.

- Kadavath et al. (2022) have found that large-parameter models are fairly well calibrated in standard benchmarks.
- However a recent OpenAI -OpenAI (2024) study found that models are fairly overconfident on tests of obscure facts:
- A Metactulus study in October 2024 found that pro human forecasters were “significantly better” than top bots on a set of 113 forecasting questions across various domains (eg. politics, economics, the rate of technological progress, etc) and that the bots displayed a positive bias (overestimating the likelihood of all events on average) ¹⁸.

3 LLM Augmentation Studies

We define AI “augmentation” as the effect of combined AI and human ability, compared to human-alone ability. These are sometimes also called “uplift” or “centaur” studies.

10



¹⁷https://assets.ctfassets.net/kftzwdyauwt9/67qJD51Aur3eIc96iOfeOP/71551c3d223cd97e591aa89567306912/o1_system_card.pdf

¹⁸<https://www.metaculus.com/notebooks/28784/aibq3results/>

3.1 The productivity effects of LLM augmentation varies widely

These RCTs compare behavior of workers with and without access to an AI system:




Figure 12: Distribution of effect sizes, from Coupé and Wu (2025)

Coupé and Wu (2025) survey 45 studies of the impact of genAI tools on productivity, with an average productivity effect of +17%, however with much variation. They also find strong signs of selection in estimates, suggesting the true effects are likely smaller.

Vaccaro, Almaatouq, and Malone (2024) survey 106 studies published between 2020 and 2023. They found:

on average, human–AI combinations performed significantly worse than the best of humans or AI alone (Hedges' $g = -0.23$; 95% confidence interval, -0.39 to -0.07). Second, we found performance losses in tasks that involved making decisions and significantly greater gains in tasks that involved creating

Domain	Reference(s)	Productivity	Quote
Software engineering	Cui et al. (2024)		
	M. Hoffmann et al. (2025)		
	Gao (2024)	+9% pull requests/day	“Accenture developers saw an 8.69% increase in pull requests ... a 15% increase to the pull request merge rate ... an 84% increase in successful builds ... developers accepted around 30% of GitHub Copilot’s suggestions.”
	Peng et al. (2023)	+126% tasks/minute	“The treatment group, with access to the AI pair programmer, completed the task 55.8% faster than the control group.” (implies +126% tasks/minute)
	Becker et al. (2025)	-16% tasks/minute	“Allowing AI actually increases completion time by 19%—AI tooling slowed developers down.” (implies -16% tasks/minute)
Consultants	Dell’Acqua et al. (2023)		
Writing tasks	Noy and Zhang (2023)		
Call centre workers	Brynjolfsson, Li, and Raymond (2023)	+15% issues/hour	“Access to AI assistance increases worker productivity, as measured by issues resolved per hour, by 15% on average, with substantial heterogeneity across workers.”
Translation	Merali (2024)	+24% translations/minute	“With access to any AI model [time taken] was reduced ... [by] 31.1%.” (=+24% translations/minute)
Entrepreneurs	Ott et al. (2024)	no statistically significant effect	
Legal tasks	Choi, Monahan, and Schwarcz (2023)	+12%–+32% tasks/minute	
Radiology	Agarwal et al. (2024)	+0.2%	38

content. Finally, when humans outperformed AI alone, we found performance gains in the combination, but when AI outperformed humans alone, we found losses.

3.2 There is some evidence of larger effects on lower-productivity workers

The following RCTs find that augmentation has a larger relative effect on workers with lower pre-experimental productivity levels:

- Brynjolfsson, Li, and Raymond (2023): Productivity increased “by 14% on average, including a 34% improvement for novice and low-skilled workers”
- Noy and Zhang (2023): “there is a correlation of 0.49 between a control participant’s average grade on the first task and their average grade on the second task. In the treatment group, initial inequalities are half-erased by the treatment: the correlation between first-task and second-task grades is only 0.25”
- Choi, Monahan, and Schwarcz (2023)
- Peng et al. (2023)
 - Magnitude: “The treatment group, with access to the AI pair programmer, completed the task 55.8% faster than the control group.”

3.3 There are few benchmarks of augmentation

- Haupt and Brynjolfsson (2025) argues for creating standardized benchmarks of augmentation/uplift they refer to as “centaur” evaluations.

3.4 Human + AI teams do not reliably outperform the better of human or AI performance individually

- Vaccaro, Almaatouq, and Malone (2024) across 16 decision-making tasks hybrid pairs under-performed the

top human-or-AI agent in 70 % of cases, though they sometimes added creative diversity.

- Agarwal et al. (2024) in a radiology-diagnosis experiment, human–AI collaboration improved accuracy by only 0.2 pp over AI-only and did **not** outperform selective automation.

3.5 Humans often misperceive the abilities of LLMs.

- Dreyfuss and Raux (2025) find “projection of human difficulty onto AI predictably distorts subjects’ beliefs”.
- Vafa, Rambachan, and Mullainathan (2024) say “We collect a dataset of 19K examples of how humans make generalizations across 79 tasks from the MMLU and BIG-Bench benchmarks ... based on short interactions, humans reach overconfident conclusions about the capabilities of the largest models.”
- Agarwal et al. (2024) find that radiologists do not use AI data efficiently.
- Becker et al. (2025) find that experienced software engineers inaccurately believe that they have experienced productivity speedups when using AI tools.

However there are limits on what these studies can tell us.

They measure short-run outcomes, but in most jobs the performance cannot be easily measured with short-run outcomes. E.g. it is challenging to interpret experiments which measure the effect of LLMs on office workers’ measurable activities (emails, pull requests).

The effects of AI are highly heterogenous across tasks and occupations.

It is difficult to extrapolate the results into the future given the advance in model capabilities (Merali (2024) attempts to fit model scale to productivity.)

These experiments often do not measure the ability of models to perform tasks autonomously.

These experiments do not incorporate general equilibrium effects.

A similar discussion of the limitations of augmentation experiments is made in a recent [Andrey Fradkin blog post](#).

4 Implications for LLM Economic Impact

4.1 The price of inference for a fixed level of capability has been falling at 10X/year or faster.

Inference price across select benchmarks, 2022–24

Source: Epoch AI, 2025; Artificial Analysis, 2025 | Chart: 2025 AI Index report





Figure 13: From Maslej et al. (2025)

- Maslej et al. (2025) estimate that inference costs, for a given level of intelligence, fell over 2022-2024 at annual rates exceeding 10X/year.
- [Artificial Analysis](#) track per-token price trends since 2022: within each tier of model intelligence, the price typically halves every few months.

11






Figure 14: From Cottier et al. (2025)

- Cottier et al. (2025) find price declines over 2023-2025 from between 9X/year (for gpt-3.5 ability on MMLU), up to 900X/year (for gpt-4o ability on GPQA).

4.2 The cost of inference for frontier capabilities has remained roughly stable.

The retail price of intelligence is determined by the compute requirement, the cost of compute, and the markup. Also note that the per-token price can be very different from the per-request price for reasoning models.

- [Artificial Analysis](#) show that the per-token price of frontier intelligence over time shows no clear trend (the chart above). Note that this chart does not include the current highest-tier models, as of July 16 Grok4 has a price of \$6/M tokens.
- The price of compute has been relatively stable over 2023-2025, with A100 priced at around \$2/hour (see [Lambda Cloud pricing](#) [Lambda Labs \(2025\)](#)).

4.3 Inference costs of serving a model are now roughly proportional to the initial training costs.

There is a tradeoff between training cost and inference cost through model size. Larger models will have higher inference costs.

- Erdil (2024) concludes AI labs should spend “comparable resources in training and running inference”.

4.4 LLMs are almost always cheaper than humans

METR (2024) compare the cost of humans and LLMs to do a variety of tasks, they say “when agents *can* do a task, they can often do so at a fraction of the cost of humans.”




Figure 6: When agents can do a task, they can often do so at a fraction of the cost of humans. The diminishing marginal improvements in agent performance as a function of cost are evident here, contrasting with the log-linear relationship observed in our human baseliners.

Figure 15: From METR (2024)

Similarly METR (2025) finds:




Figure 16: Agent Performance on HCAST & RE-Bench by Cost

Exceptions:

- Computer transcription was for a period more accurate but less cost-effective than humans.
- Reasoning tasks with high test-time expenditure ([ARC-AGI](#))
- [Svanberg et al. \(2024\)](#) argue that “only 23% of worker compensation “exposed” to AI computer vision would be cost-effective for firms to automate because of the large upfront costs of AI systems.” However this is almost entirely fixed costs (data collection, fine tuning, scaffolding), not marginal costs.

4.5 The economic impact of LLMs has been modest so far

Around 1/2 of US adults regularly use an LLM-based chatbot. The productivity impact is perhaps 0.1-1% of GDP:

- [Bick, Blandin, and Deming \(2024\)](#): 0.1-0.9% aggregate productivity boost through time-savings on work tasks.
- [Humlum and Vestergaard \(2025\)](#): 3% average self-reported time saving among users.

- Collis and Brynjolfsson (2025): genAI app users report valuations (Willingness-to-Accept) equivalent to around 0.3% of GDP (\$98 per user, \$97B total).

Probably the biggest economic impact is on home production.

Working on a paper documenting ChatGPT use. Note that Anthropic and Microsoft papers only focus on use at work tasks.

Enterprise impacts are still small.

Most enterprises have adopted AI in some form, but they typically report small effects on revenue or costs.


Maslej et al. (2025) say:

“Most companies that report financial impacts from using AI within a business function estimate the benefits as being at low levels. 49% of respondents whose organizations use AI in service operations report cost savings, followed by supply chain management (43%) and software engineering (41%), but most of them report cost savings of less than 10%.”

Most growth in impact is likely due to capabilities rather than diffusion. Thought experiment: if we had frozen capabilities at gpt-3.5, gpt-4.0, what would be the impact today?

We can observe the causal effect of model improvement on existing cohorts with AB tests, it is harder to attribute user acquisition.

12



4.6 We may not be able to direct technical change

Many economists have said we should shift the focus of our AI research:

Brynjolfsson (2023)

“an excessive focus on developing and deploying [human-level AI] can lead us into a trap ... when AI is focused on augmenting humans rather than

mimicking them, then humans retain the power to insist on a share of the value created”

Acemoglu (2021):

“our current trajectory automates work to an excessive degree while refusing to invest in human productivity

This may not be possible: augmenting AI and automating AI may be the same AI.

If there is a single latent factor of model ability, then work to train an augmenting model looks very similar to work to train an automating model. The ability to significantly steer AI capabilities seems limited.

4.7 A univariate model of machine and human intelligence will be a poor fit




Figure 17: from Aschenbrenner (2024)

We sometimes describe AI capability by human equivalent, e.g. years of education (Aschenbrenner (2024)).

What would a two-factor model look like?

- LLMs have PhD-level knowledge and reasoning.
- LLMs have preschool-level tactics.

4.8 Wages are unlikely to be pinned down by the price of compute

Some models feature “equilibrium task-sharing”: humans perform tasks that computers are able to do, but it is too expensive for computers:

- Trammell and Korinek (2023)
- Ide and Talamas (2024)
- Noah Smith (2024) on comparative advantage and wages.
- Acemoglu (2024)
- Jones (2025) assumes that all tasks which an AI can do will be done by an AI in equilibrium (assumption 1, p8). However he also expects that significant progress through a reduction in the price of AI tasks.

We argued above that equilibrium task-sharing has historically been true only for a small share of tasks.

Equilibrium task-sharing would require either:

1. A dramatic increase in the price to perform a task:
 - An increase in the compute required for new tasks – the price of frontier inference has been steady.
 - An increase in compute costs – GPU rental prices have been steady, and the purchase/rent ratio has been stable.
 - An increase in margins.
2. A dramatic drop in human wages.

5 Appendix / Offcuts

Jason Wei blog post “*The ease of training AI to solve a task is proportional to how verifiable the task is. All tasks that are possible to solve and easy to verify will be solved by AI*”. Five specific factors:

1. Objective truth: everyone agrees what good solutions are
2. Fast to verify: any given solution can be verified in a few seconds
3. Scalable to verify: many solutions can be verified simultaneously
4. Low noise: verification is as tightly correlated to the solution quality as possible
5. Continuous reward: it’s easy to rank the goodness of many solutions for a single problem

References

- Acemoglu, Daron. 2021. “AI’s Future Doesn’t Have to Be Dystopian.” *Boston Review* 20.
- . 2024. “The Simple Macroeconomics of AI.” National Bureau of Economic Research.
- Agarwal, Nikhil, Alex Moehring, Pranav Rajpurkar, and Tobias Salz. 2024. “Combining Human Expertise with Artificial Intelligence: Experimental Evidence from Radiology.” MIT Department of Economics. <https://economics.mit.edu/sites/default/files/inline-files/agarwal-et-al-diagnostic-ai%20%285%29.pdf>.
- Aschenbrenner, Leopold. 2024. “Situational Awareness.” *The Decade Ahead. Situational-Awareness. Ai*.
- Backlund, Axel, and Lukas Petersson. 2025. “Vending-Bench: A Benchmark for Long-Term Coherence of Autonomous Agents.” *arXiv Preprint arXiv:2502.15840*. <https://arxiv.org/abs/2502.15840>.
- Bai, Yushi, Shangqing Tu, Jiajie Zhang, Hao Peng, Xiaozhi Wang, Xin Lv, Shulin Cao, et al. 2024. “LongBench V2: Towards Deeper Understanding and Reasoning on

- Realistic Long-Context Multitasks.” *arXiv Preprint arXiv:2412.15204*. <https://arxiv.org/abs/2412.15204>.
- Balachandran, Vidhisha, Jingya Chen, Lingjiao Chen, Shivam Garg, Neel Joshi, Yash Lara, John Langford, et al. 2025. “Inference-Time Scaling for Complex Tasks: Where We Stand and What Lies Ahead.” *arXiv Preprint arXiv:2504.00294*.
- Becker, Joel, Nate Rush, Beth Barnes, and David Rein. 2025. “Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity.” *arXiv Preprint arXiv:2507.09089*. <https://arxiv.org/abs/2507.09089>.
- Beeching, Edward, Lewis Tunstall, and Sasha Rush. n.d. “Scaling Test-Time Compute with Open Models.” <https://huggingface.co/spaces/HuggingFaceH4/blogpost-scaling-test-time-compute>.
- Bick, Alexander, Adam Blandin, and David J Deming. 2024. “The Rapid Adoption of Generative AI.” National Bureau of Economic Research.
- Brynjolfsson, Erik. 2023. “The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence.” In *Augmented Education in the Global Age*, 103–16. Routledge.
- Brynjolfsson, Erik, Danielle Li, and Lindsey R. Raymond. 2023. “Generative AI at Work.” *arXiv Preprint arXiv:2304.11771*. <https://arxiv.org/abs/2304.11771>.
- Burnell, Richard, Heng Hao, Andrew R. A. Conway, and José Hernández-Orallo. 2023. “Revealing the Structure of Language Model Capabilities.” *arXiv Preprint arXiv:2306.10062*. <https://arxiv.org/abs/2306.10062>.
- Choi, Jonathan H., Amy Monahan, and Daniel Schwarcz. 2023. “AI Assistance in Legal Work: Evidence from a Field Experiment.” *SSRN Electronic Journal*. <https://doi.org/10.2139/ssrn.4626276>.
- Chollet, François. 2019. “On the Measure of Intelligence.” <https://arxiv.org/abs/1911.01547>.
- Choshen, Leshem, Yang Zhang, and Jacob Andreas. 2024. “A Hitchhiker’s Guide to Scaling Law Estimation.” *arXiv Preprint arXiv:2410.11840*.
- Chui, Michael, Eric Hazan, Roger Roberts, Alex Singla, and Kate Smaje. 2023. “The Economic Potential of Generative AI.” [https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-](https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-ai)

- [ai-the-next-productivity-frontier](#).
- Collis, Avinash, and Erik Brynjolfsson. 2025. “AI’s Overlooked \$97 Billion Contribution to the Economy.” *The Wall Street Journal*, August. <https://www.wsj.com/opinion/ais-overlooked-97-billion-contribution-to-the-economy-users-service-da6e8f55>.
- Cottier, Ben, Ben Snodin, David Owen, and Tom Adamczewski. 2025. “LLM Inference Prices Have Fallen Rapidly but Unequally Across Tasks.” <https://epoch.ai/data-insights/llm-inference-price-trends>.
- Coupé, Tom, and Weilun Wu. 2025. “The Impact of Generative AI on Productivity: Results of an Early Meta-Analysis.” 25/09. University of Canterbury, Department of Economics; Finance. <https://doi.org/None>.
- Cui, Yiming, Mert Demirer, Sonia Jaffe, Andreas Musolff, Sida Peng, and Tobias Salz. 2024. “AI-Augmented Software Engineering: Evidence from a Large-Scale Field Experiment.” *SSRN Electronic Journal*. <https://doi.org/10.2139/ssrn.4945566>.
- Dell’Acqua, Fabrizio, Edward McFowland, Ethan R. Mollick, Hila Lifshitz-Assaf, Katherine Kellogg, Saran Rajendran, Lisa Kraye, François Candelon, and Karim R. Lakhani. 2023. “Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality.” *Harvard Business School Working Paper*. <https://www.hbs.edu/faculty/Pages/item.aspx?num=64700>.
- Dorn, Félix, Xavier Ferrés, Elsie Jang, and Valmik Nahata. 2025. “The Early Economic Impacts of Transformative AI: A Focus on Temporal Coherence.” Research submission for the „Economics of Transformative AI” initiative, organised by Apart Research. <https://apartresearch.com/project/the-early-economic-impacts-of-transformative-ai-a-focus-on-temporal-coherence-ipql>.
- Dreyfuss, Bnaya, and Raphaël Raux. 2025. “Human Learning about AI: Projection of Difficulty Distorts Beliefs.” *arXiv Preprint arXiv:2406.05408*. <https://arxiv.org/abs/2406.05408>.
- Eloundou, Tyna, Sam Manning, Pamela Mishkin, and Daniel Rock. 2023. “Gpts Are Gpts: An Early Look at the Labor Market Impact Potential of Large Language Models.” *arXiv*

Preprint arXiv:2303.10130.

- Erdil, Ege. 2024. “Optimally Allocating Compute Between Inference and Training.” <https://epoch.ai/blog/optimally-allocating-compute-between-inference-and-training>.
- Gao, Albert. 2024. “Research: Quantifying GitHub Copilot’s Impact in the Enterprise with Accenture.” <https://github.blog/2024-05-13-research-quantifying-github-copilots-impact-in-the-enterprise-with-accenture/#from-the-lab-to-the-real-world>.
- Google. 2025. “AlphaEvolve: A Gemini-Powered Coding Agent for Designing Advanced Algorithms.” <https://deepmind.google/discover/blog/alphaevolve-a-gemini-powered-coding-agent-for-designing-advanced-algorithms/>.
- Haupt, Andreas, and Erik Brynjolfsson. 2025. “Position: AI Should Not Be an Imitation Game: Centaur Evaluations.” In *Forty-Second International Conference on Machine Learning Position Paper Track*. <https://openreview.net/forum?id=LkdH35003E>.
- Heitmann, Max, Ture Hinrichsen, David Africa, and Jonas Sandbrink. 2025. “Understanding AI Trajectories: Mapping the Limitations of Current AI Systems (24_10 Update).” AI Safety Institute. [https://cdn.prod.website-files.com/663bd486c5e4c81588db7a1d/68fb86aa2c3b1b7ea6251cc1_Understanding%20AI%20Trajectories%20\(24_10%20update\).pdf](https://cdn.prod.website-files.com/663bd486c5e4c81588db7a1d/68fb86aa2c3b1b7ea6251cc1_Understanding%20AI%20Trajectories%20(24_10%20update).pdf).
- Hendrycks, Dan, Dawn Song, Christian Szegedy, Honglak Lee, Yarin Gal, Erik Brynjolfsson, Sharon Li, et al. 2025. “A Definition of AGI.” <https://www.agidefinition.ai/paper.pdf>.
- Hoffmann, Jordan, Sebastian Borgeaud, Arthur Mensch, et al. 2022. “Training Compute-Optimal Large Language Models.” *arXiv Preprint arXiv:2203.15556*. <https://arxiv.org/abs/2203.15556>.
- Hoffmann, Manuel, Sam Boysel, Frank Nagle, Sida Peng, and Kevin Xu. 2025. “Generative AI and the Nature of Work.” *Harvard Business School Strategy Unit Working Paper*, no. 25-021: 25–021.
- Hu, Lanxiang, Mingjia Huo, Yuxuan Zhang, Haoyang Yu, Eric P. Xing, Ion Stoica, Tajana Rosing, Haojian Jin, and Hao Zhang. 2025. “Lmgame-Bench: How Good Are LLMs at Playing Games?” <https://arxiv.org/abs/2505.15146>.
- Huber, Thomas, and Christina Niklaus. 2025. “LLMs Meet

- Bloom’s Taxonomy: A Cognitive View on Large Language Model Evaluations.” In *Proceedings of the 31st International Conference on Computational Linguistics*, 5211–46. Abu Dhabi, UAE: Association for Computational Linguistics. <https://aclanthology.org/2025.coling-main.350/>.
- Humlum, Anders, and Emilie Vestergaard. 2025. “Large Language Models, Small Labor Market Effects.” National Bureau of Economic Research.
- Ide, Enrique, and Eduard Talamas. 2024. “Artificial Intelligence in the Knowledge Economy.” <https://arxiv.org/abs/2312.05481>.
- Ilić, David, and Gilles E. Gignac. 2024. “Evidence of Interrelated Cognitive-Like Capabilities in Large Language Models: Indications of Artificial General Intelligence or Achievement?” *Intelligence* 106 (September): 101858. <https://doi.org/10.1016/j.intell.2024.101858>.
- Jones, Benjamin F. 2025. “Artificial Intelligence in Research and Development.” Working Paper w34312. National Bureau of Economic Research. <https://doi.org/10.3386/w34312>.
- Kadavath, Saurav, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, et al. 2022. “Language Models (Mostly) Know What They Know.” <https://arxiv.org/abs/2207.05221>.
- Kaplan, Jared, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. “Scaling Laws for Neural Language Models.” <https://arxiv.org/abs/2001.08361>.
- Kipnis, Alex, Konstantinos Voudouris, Luca M. Schulze Buschoff, and Eric Schulz. 2025. “Metabench - a Sparse Benchmark of Reasoning and Knowledge in Large Language Models.” In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=4T33izzFpK>.
- Kwa, Timothy, Xia Liu, Yu Yuan, Tanish Singh, et al. 2025. “Measuring AI Ability to Complete Long Tasks.” *arXiv Preprint arXiv:2503.14499*. <https://arxiv.org/abs/2503.14499>.
- Lambda Labs. 2025. “Lambda Cloud GPU Pricing.” <https://lambdalabs.com/service/gpu-cloud/pricing>.

- Legg, Shane, and Marcus Hutter. 2007a. “A Collection of Definitions of Intelligence.” *Frontiers in Artificial Intelligence and Applications* (Vol. 157), 17–24. <https://arxiv.org/abs/0706.3639>.
- . 2007b. “Universal Intelligence: A Definition of Machine Intelligence.” *Minds and Machines* 17 (4): 391–444. <https://doi.org/10.1007/s11023-007-9079-x>.
- Li, Xinyuan, Murong Xu, Wenbiao Tao, Hanlun Zhu, Yike Zhao, Jipeng Zhang, and Yunshi Lan. 2025. “RIDE: Difficulty Evolving Perturbation with Item Response Theory for Mathematical Reasoning.” <https://arxiv.org/abs/2511.04120>.
- Maimon, Aviya, Amir DN Cohen, Gal Vishne, Shauli Ravfogel, and Reut Tsarfaty. 2025. “IQ Test for LLMs.” *arXiv Preprint arXiv:2507.20208*.
- Maslej, Nestor, Loredana Fattorini, Raymond Perrault, Yolanda Gil, Vanessa Parli, Njenga Kariuki, Emily Capstick, et al. 2025. “Artificial Intelligence Index Report 2025.” *arXiv Preprint arXiv:2504.07139*.
- Merali, Ali. 2024. “Scaling Laws for Economic Productivity: Experimental Evidence in LLM-Assisted Translation.” *arXiv Preprint arXiv:2409.02391*. <https://arxiv.org/abs/2409.02391>.
- METR. 2024. “An Update on Our General Capability Evaluations.” <https://metr.org/blog/2024-08-06-update-on-evaluations/>.
- . 2025. “Details about METR’s Evaluation of OpenAI GPT-5.” <https://evaluations.metr.org/gpt-5-report/>.
- Min, Sewon, Kalpesh Krishna, Xinxu Lyu, et al. 2023. “FACTScore: Fine-Grained Atomic Evaluation of Factual Precision in Long Form Text Generation.” *arXiv Preprint arXiv:2305.14251*. <https://arxiv.org/abs/2305.14251>.
- Mirzadeh, Iman, Keivan Alizadeh, Hooman Shahrokhi, Oncel Tuzel, Samy Bengio, and Mehrdad Farajtabar. 2024. “GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models.” *arXiv Preprint arXiv:2410.05229*. <https://arxiv.org/abs/2410.05229>.
- Miserendino, Samuel, Michele Wang, Tejal Patwardhan, and Johannes Heidecke. 2025. “SWE-Lancer: Can Frontier LLMs Earn \$1 Million from Real-World Freelance Software

- Engineering?” <https://arxiv.org/abs/2502.12115>.
- Morris, Meredith Ringel, Jascha Sohl-Dickstein, Noah Fiedel, Tris Warkentin, Allan Dafoe, Aleksandra Faust, Clement Farabet, and Shane Legg. 2023. “Levels of AGI for Operationalizing Progress on the Path to AGI.” *arXiv Preprint arXiv:2311.02462*.
- Noy, Shakked, and Whitney Zhang. 2023. “Experimental Evidence on the Productivity Effects of Generative Artificial Intelligence.” *Science* 381 (6654): 187–92. <https://doi.org/10.1126/science.adh2586>.
- OECD. 2025. “Introducing the OECD AI Capability Indicators.” Paris: OECD Publishing. <https://doi.org/10.1787/be745f04-en>.
- OpenAI. 2024. “Introducing SimpleQA.” <https://openai.com/index/introducing-simpleqa/>.
- Ord, Toby. 2025. “The Half-Life of AI Success: A Model of Performance Decay.” <https://www.tobyord.com/writing/half-life>.
- Otis, Christopher et al. 2024. “AI-Augmented Entrepreneurship: Evidence from a Field Experiment.” *SSRN Electronic Journal*. <https://doi.org/10.2139/ssrn.4671369>.
- Owen, David. 2024. “How Predictable Is Language-Model Benchmark Performance?” *arXiv Preprint arXiv:2401.04757*. <https://arxiv.org/abs/2401.04757>.
- Payne, Kenneth, and Baptiste Alloui-Cros. 2025. “Strategic Intelligence in Large Language Models: Evidence from Evolutionary Game Theory.” *arXiv Preprint arXiv:2507.02618*. <https://arxiv.org/abs/2507.02618>.
- Peng, Sida, Eirini Kalliamvakou, Peter Cihon, and Mert Demirer. 2023. “The Impact of AI on Developer Productivity: Evidence from GitHub Copilot.” <https://arxiv.org/abs/2302.06590>.
- Phan, Long, Mantas Mazeika, Andy Zou, and Dan Hendrycks. 2025. “TextQuests: How Good Are LLMs at Text-Based Video Games?” *arXiv Preprint arXiv:2507.23701*. <https://arxiv.org/abs/2507.23701>.
- Polo, Felipe Maia, Seamus Somerstep, Leshem Choshen, Yuekai Sun, and Mikhail Yurochkin. 2025. “Sloth: Scaling Laws for LLM Skills to Predict Multi-Benchmark Performance Across Families.” <https://arxiv.org/abs/2412.06540>.
- Renze, Matthew, and Erhan Guven. 2024. “The Effect of

- Sampling Temperature on Problem Solving in Large Language Models.” *arXiv Preprint arXiv:2402.05201*. <https://arxiv.org/abs/2402.05201>.
- Ruan, Yizhong, Chris J. Maddison, and Tatsunori B. Hashimoto. 2024. “Observational Scaling Laws and the Predictability of Language-Model Performance.” *arXiv Preprint arXiv:2405.10938*. <https://arxiv.org/abs/2405.10938>.
- Song, Yifan, Guoyin Wang, Sujian Li, and Bill Yuchen Lin. 2024. “The Good, the Bad, and the Greedy: Evaluation of LLMs Should Not Ignore Non-Determinism.” *arXiv Preprint arXiv:2407.10457*. <https://arxiv.org/abs/2407.10457>.
- Spangher, Lewis, Tao Li, William F. Arnold, et al. 2024. “Towards a Generalized Performance Benchmark for LLM Capabilities.” *arXiv Preprint arXiv:2410.22368*. <https://arxiv.org/abs/2410.22368>.
- Sutton, Richard. 2019. “The Bitter Lesson.” *Incomplete Ideas (Blog)* 13 (1): 38. <http://www.incompleteideas.net/IncIdeas/BitterLesson.html>.
- Toner, Helen. 2025. “Taking Jaggedness Seriously.” Substack — Rising Tide.
- Trammell, Philip, and Anton Korinek. 2023. “Economic Growth Under Transformative AI.” National Bureau of Economic Research.
- Vaccaro, Michelle, Abdullah Almaatouq, and Thomas Malone. 2024. “When Combinations of Humans and AI Are Useful: A Systematic Review and Meta-Analysis.” *Nature Human Behaviour* 8: 2293–303. <https://doi.org/10.1038/s41562-024-02024-1>.
- Vafa, Keyon, Justin Y. Chen, Ashesh Rambachan, Jon Kleinberg, and Sendhil Mullainathan. 2024. “Evaluating the World Model Implicit in a Generative Model.” In *Advances in Neural Information Processing Systems*, edited by A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, 37:26941–75. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2024/file/2f6a6317bada76b26a4f61bb70a7db59-Paper-Conference.pdf.
- Vafa, Keyon, Ashesh Rambachan, and Sendhil Mullainathan. 2024. “Do Large Language Models Perform the Way Peo-

- ple Expect? Measuring the Human Generalization Function.” *arXiv Preprint arXiv:2406.01382*. <https://arxiv.org/abs/2406.01382>.
- Valmeekam, Karthik, Kaya Stechly, and Subbarao Kambhampati. 2024. “LLMs Still Can’t Plan; Can LRMs? A Preliminary Evaluation of OpenAI’s O1 on Plan-Bench.” *arXiv Preprint arXiv:2409.13373*. <https://arxiv.org/abs/2409.13373>.
- Wang, Xuezhi, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. “Self-Consistency Improves Chain of Thought Reasoning in Language Models.” <https://arxiv.org/abs/2203.11171>.
- Wei, Jerry, Chengrun Yang, Xinying Song, et al. 2024. “Long-Form Factuality in Large Language Models.” *arXiv Preprint arXiv:2403.18802*. <https://arxiv.org/abs/2403.18802>.
- Werra, Leandro von. 2025. “The Jagged AI Frontier Is a Data Frontier.” 2025. <https://huggingface.co/spaces/lvwerra/jagged-data-frontier>.
- Wu, Yangzhen, Zhiqing Sun, Shanda Li, Sean Welleck, and Yiming Yang. 2025. “Inference Scaling Laws: An Empirical Analysis of Compute-Optimal Inference for Problem-Solving with Language Models.” <https://arxiv.org/abs/2408.00724>.
- You, Josh. 2025. “How Far Can Reasoning Models Scale?” <https://epoch.ai/gradient-updates/how-far-can-reasoning-models-scale>.
- Zhou, Lexin, Lorenzo Pacchiardi, Fernando Martínez-Plumed, Katherine M. Collins, Yael Moros-Daval, Seraphina Zhang, Qinlin Zhao, et al. 2025. “General Scales Unlock AI Evaluation with Explanatory and Predictive Power.” <https://arxiv.org/abs/2503.06378>.
- Zhou, Lucy, Winfried Schellaert, Fernando Martínez-Plumed, et al. 2024. “Larger and More Instructable Language Models Become Less Reliable.” *Nature* 634: 61–69. <https://doi.org/10.1038/s41586-024-07930-y>.